

Alright, here's our situation."

We're sitting on a dataset of over **8,000 used car listings** from the Indian market. This is our inventory intelligence — it tells us what's been sold, at what price, and in what condition.

What we have on each car:

- The car's name, make, and model year
- The **selling price** — our core number
- How many kilometers it's been driven
- Fuel type (Petrol, Diesel, CNG, LPG)
- Whether it was sold by an **individual or a dealer**
- Transmission type (Manual or Automatic)
- Ownership history — first owner, second, third, etc.
- Technical specs: mileage, engine size, power, torque, seats

What problem am I trying to solve?

Right now, when a car comes in, my team is **guessing at the right price**. We look at the name, we look at the year, and we throw a number at it. We're either leaving money on the table or scaring away buyers.

I want to know: **What actually drives the price of a used car?** Is it the kilometers driven? The fuel type? Whether it's a first owner? I need to stop guessing.

I also have a **data quality problem** — some listings are missing mileage, engine specs, even seat counts. That's hurting our ability to make decisions confidently.

What I want from this data:

1. A reliable way to **predict the fair selling price** of any car we bring in
2. To understand **which factors matter most** to price — so my sales team knows what to highlight
3. To clean up the gaps so every listing tells a complete story